

CA4 - Stellar object classification

Import libraries

```
# import libraries for plotting and data manipulation
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# import classifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier

# import model selection and preprocessing tools
from sklearn.decomposition import PCA
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import StratifiedKFold, GridSearchCV,
train_test_split
from sklearn.metrics import confusion_matrix, classification_report,
f1_score
```

Reading the data

```
# load training and test data
training_data = pd.read_csv('data/train.csv')
test_data = pd.read_csv('data/test.csv')
```

Data exploration

Look at the shape of the data

```
# check the shape of the data
training_data.shape, test_data.shape

((80000, 18), (20000, 17))
```

We can see by the shape of the data, that the training and test data is split 80/20

Overview of the data

```
training_data.head()
```

	obj_ID	alpha	delta	u	g
r \					

0	1.237661e+18	135.689107	32.494632	23.87882	22.27530	20.39501
1	1.237665e+18	144.826101	31.274185	24.77759	22.83188	22.58444
2	1.237661e+18	142.188790	35.582444	25.26307	22.66389	20.60976
3	1.237663e+18	338.741038	-0.402828	22.13682	23.77656	21.61162
4	1.237680e+18	345.282593	21.183866	19.43718	17.58028	16.49747

	i	z	run_ID	rerun_ID	cam_col	field_ID
spec_obj_ID \						
0	19.16573	18.79371	3606	301	2	79
6.543777e+18						
1	21.16812	21.61427	4518	301	5	119
1.176014e+19						
2	19.34857	18.94827	3606	301	2	120
5.152200e+18						
3	20.50454	19.25010	4192	301	3	214
1.030107e+19						
4	15.97711	15.54461	8102	301	3	137
6.891865e+18						

	class	redshift	plate	MJD	fiber_ID
0	GALAXY	0.634794	5812	56354	171
1	GALAXY	0.779136	10445	58158	427
2	GALAXY	0.644195	4576	55592	299
3	GALAXY	0.932346	9149	58039	775
4	GALAXY	0.116123	6121	56187	842

The data has 17 columns, where the class column is the class of the stellar object, our target variable y.

Visualize feature distributions using violinplots

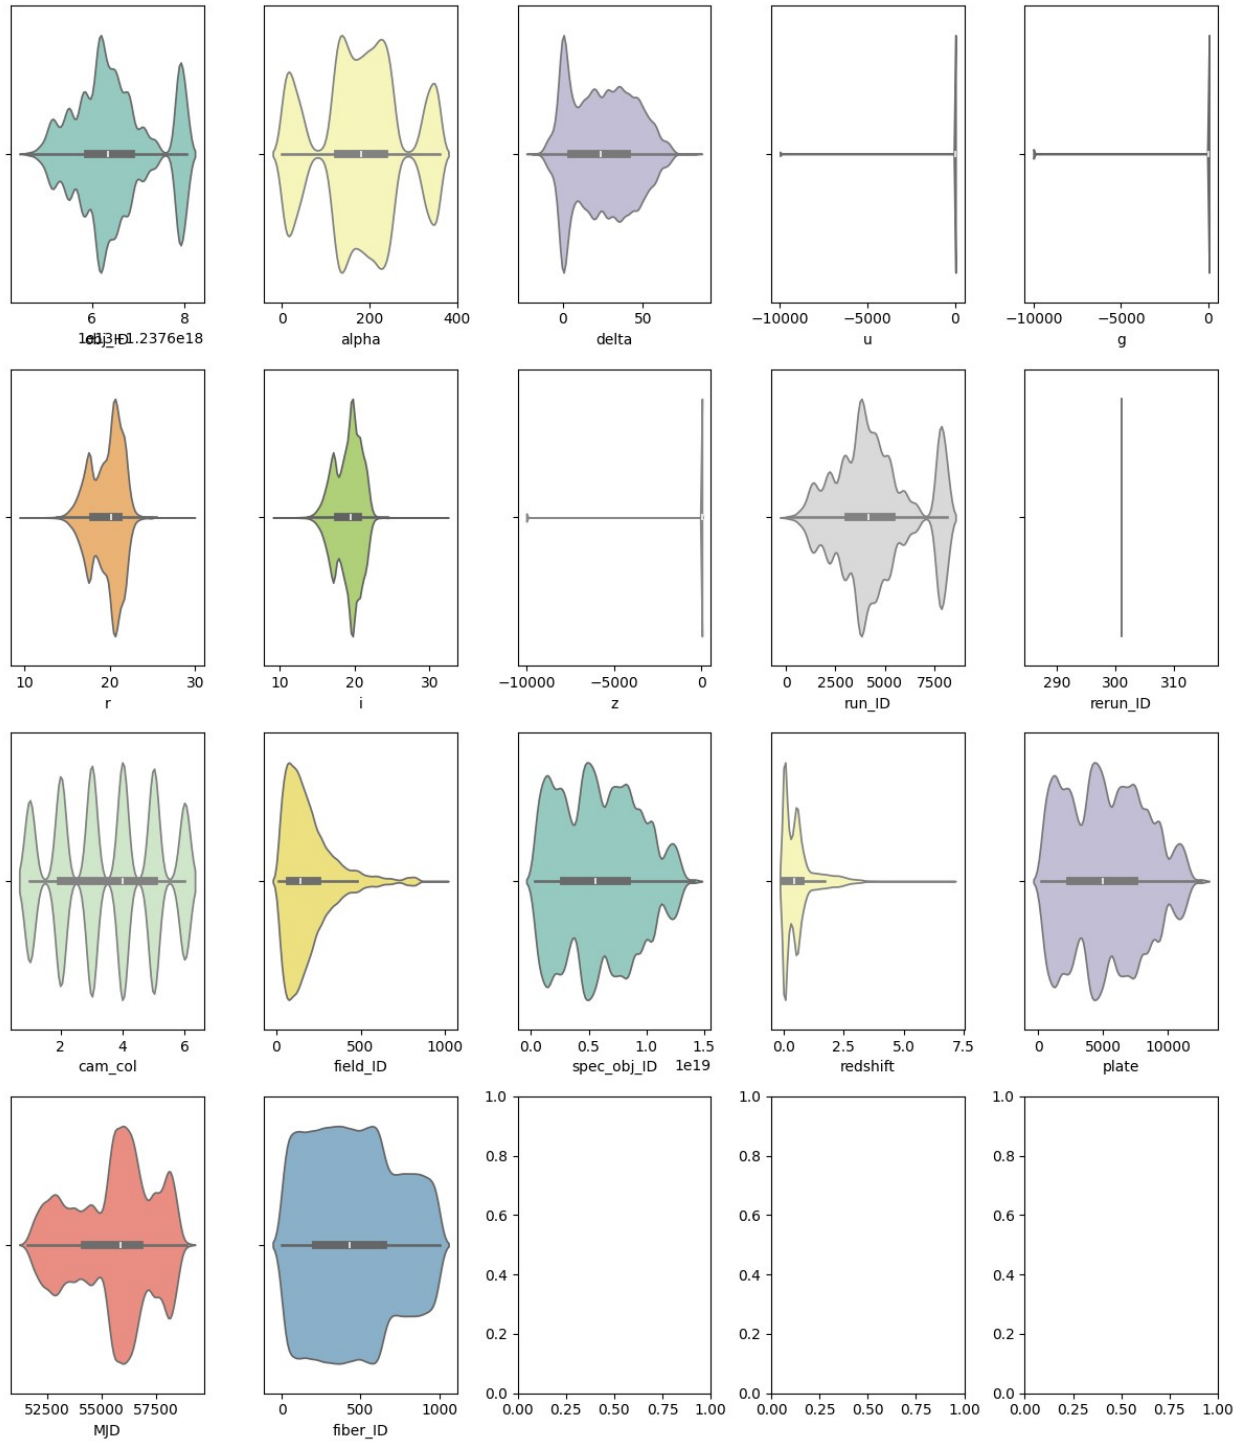
```
# visualize feature distributions
fig, axes = plt.subplots(4,5, figsize = (12,14))
axes = axes.flatten()

columns = training_data.columns.to_list()
columns = [col for i, col in enumerate(columns) if i != 13]

colors = sns.color_palette('Set3', n_colors=len(columns))

for i, col in enumerate(columns):
    sns.violinplot(x=col, data=training_data, ax=axes[i],
color=colors[i])

plt.tight_layout()
plt.show()
```



Perform descriptive statistics

```
training_data.describe()
```

	obj_ID	alpha	delta	u
g \				
count	8.000000e+04	80000.000000	80000.000000	79638.000000

```

80000.000000
mean    1.237665e+18    177.579220    24.132590    21.961115
20.507677
std      8.424878e+12    96.409584    19.650113    35.581856
35.483302
min      1.237646e+18    0.005528    -18.785328    -9999.000000 -
9999.000000
25%      1.237659e+18    127.643892    5.170723    20.353990
18.963188
50%      1.237663e+18    180.761747    23.603480    22.187965
21.101015
75%      1.237668e+18    233.815698    39.904905    23.698457
22.125007
max      1.237681e+18    359.999615    83.000519    32.781390
31.602240

```

```

              r              i              z      run_ID
rerun_ID \
count  80000.000000  80000.000000  80000.000000  80000.000000
80000.0
mean      19.647426    19.085051    18.644640    4477.876713
301.0
std        1.855636    1.757630    35.462189    1961.579187
0.0
min        9.822070    9.469903   -9999.000000    109.000000
301.0
25%       18.135523    17.732600    17.457503    3185.000000
301.0
50%       20.127550    19.405635    19.004420    4188.000000
301.0
75%       21.047242    20.401857    19.923013    5326.000000
301.0
max       29.571860    32.141470    29.383740    8162.000000
301.0

```

```

      cam_col    field_ID    spec_obj_ID    redshift
plate \
count  80000.000000  80000.000000  8.000000e+04  80000.000000
80000.000000
mean      3.511388    185.663050  5.783094e+18    0.577219
5136.309963
std        1.589033    148.433559  3.327780e+18    0.731597
2955.646539
min        1.000000    11.000000  2.995801e+17    -0.009971
266.000000
25%        2.000000    82.000000  2.841535e+18    0.055160
2523.750000
50%        4.000000    146.000000  5.606066e+18    0.425051
4979.000000

```

75%	5.000000	240.000000	8.331746e+18	0.704777
7400.000000				
max	6.000000	989.000000	1.412694e+19	7.011245
12547.000000				

	MJD	fiber_ID
count	80000.000000	80000.000000
mean	55587.191400	448.887437
std	1809.144282	272.328619
min	51608.000000	1.000000
25%	54233.000000	221.000000
50%	55868.000000	432.000000
75%	56777.000000	644.000000
max	58932.000000	1000.000000

Data cleaning

Look for missing values

```
# identify rows with missing values
NaN = training_data.isna().sum()
print(NaN)

print(f'Missing values in the training data: {NaN[3]}')
```

```
obj_ID      0
alpha       0
delta       0
u           362
g           0
r           0
i           0
z           0
run_ID      0
rerun_ID    0
cam_col     0
field_ID    0
spec_obj_ID 0
class       0
redshift    0
plate       0
MJD         0
fiber_ID    0
dtype: int64
Missing values in the training data: 362
```

```
C:\Users\anekl\AppData\Local\Temp\ipykernel_8884\2165697918.py:5:
FutureWarning: Series.__getitem__ treating keys as positions is
deprecated. In a future version, integer keys will always be treated
as labels (consistent with DataFrame behavior). To access a value by
```

```
position, use `ser.iloc[pos]`  
print(f'Missing values in the training data: {NaN[3]}')
```

Remove missing values from the training data

```
# remove rows with missing values  
training_data_clean = training_data.dropna()  
  
# check if there are any missing values left  
training_data_clean.isna().sum()  
  
print(f'Missing values left in the training data:  
{training_data_clean.isna().sum().sum()}')
```

Missing values left in the training data: 0

Remove excessive features from the data set

```
training_data_clean = training_data_clean.drop(columns=['obj_ID',  
'run_ID', 'rerun_ID', 'cam_col', 'field_ID', 'spec_obj_ID', 'MJD',  
'fiber_ID', 'plate', 'fiber_ID'])  
training_data_clean.head()
```

	alpha	delta	u	g	r	i
z \						
0	135.689107	32.494632	23.87882	22.27530	20.39501	19.16573
18.79371						
1	144.826101	31.274185	24.77759	22.83188	22.58444	21.16812
21.61427						
2	142.188790	35.582444	25.26307	22.66389	20.60976	19.34857
18.94827						
3	338.741038	-0.402828	22.13682	23.77656	21.61162	20.50454
19.25010						
4	345.282593	21.183866	19.43718	17.58028	16.49747	15.97711
15.54461						

	class	redshift
0	GALAXY	0.634794
1	GALAXY	0.779136
2	GALAXY	0.644195
3	GALAXY	0.932346
4	GALAXY	0.116123

```
# remove excessive features from the test data  
test_data_clean = test_data.drop(columns=['obj_ID', 'run_ID',  
'rerun_ID', 'cam_col', 'field_ID', 'spec_obj_ID', 'MJD', 'fiber_ID',  
'plate', 'fiber_ID'])  
test_data_clean.head()
```

	alpha	delta	u	g	r	i
z \						

0	16.956890	3.646130	23.33542	21.95143	20.48149	19.60300
1	240.063240	6.134131	17.86033	16.79228	16.43001	16.30923
2	30.887222	1.188710	18.18911	16.89469	16.42161	16.24627
3	247.594401	10.887780	24.99961	21.71203	21.47148	21.30532
4	18.896451	-5.261330	23.76648	21.79737	20.69543	20.23403

	redshift
0	0.506237
1	0.000345
2	0.000004
3	-0.000291
4	-0.000136

Visualize feature distributions after removing excessive features


```
# visualize feature distributions
fig, axes = plt.subplots(2,4, figsize = (10,5))
axes = axes.flatten()

columns = training_data_clean.columns.to_list()
columns = [col for i, col in enumerate(columns) if i != 7]

colors = sns.color_palette('Set3', n_colors=len(columns))

for i, col in enumerate(columns):
    sns.boxplot(x=col, data=training_data_clean, ax=axes[i],
color=colors[i])

plt.tight_layout()
plt.show()
```



Visualization using violinplots

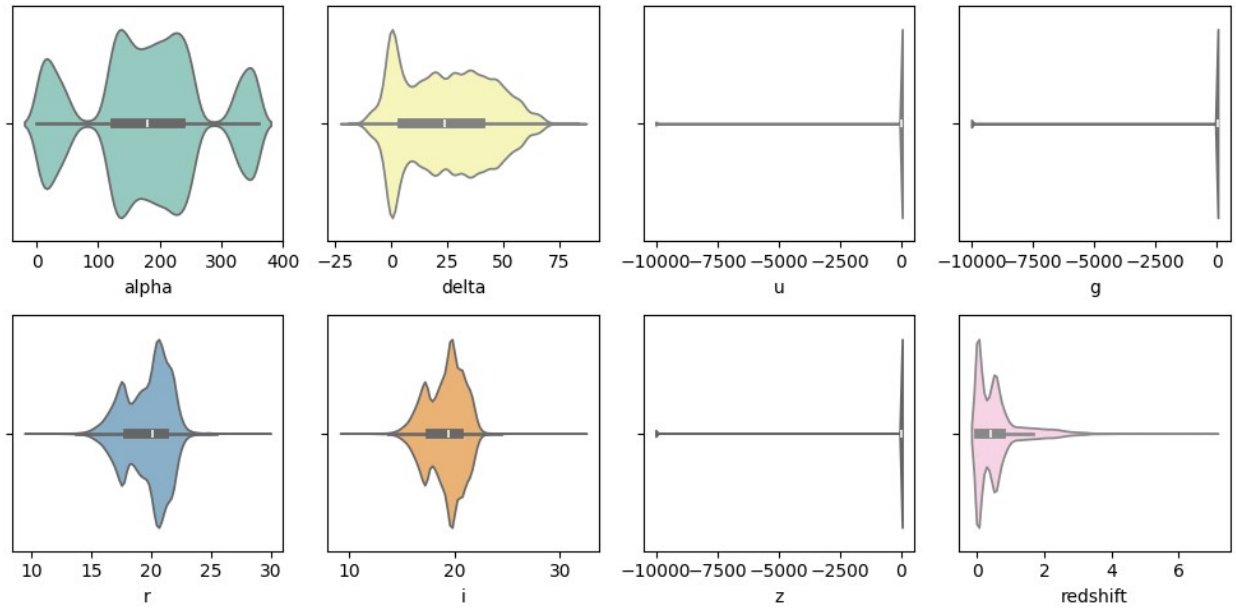
```
# visualize feature distributions
fig, axes = plt.subplots(2,4, figsize = (10,5))
axes = axes.flatten()

columns = training_data_clean.columns.to_list()
columns = [col for i, col in enumerate(columns) if i != 7]

colors = sns.color_palette('Set3', n_colors=len(columns))

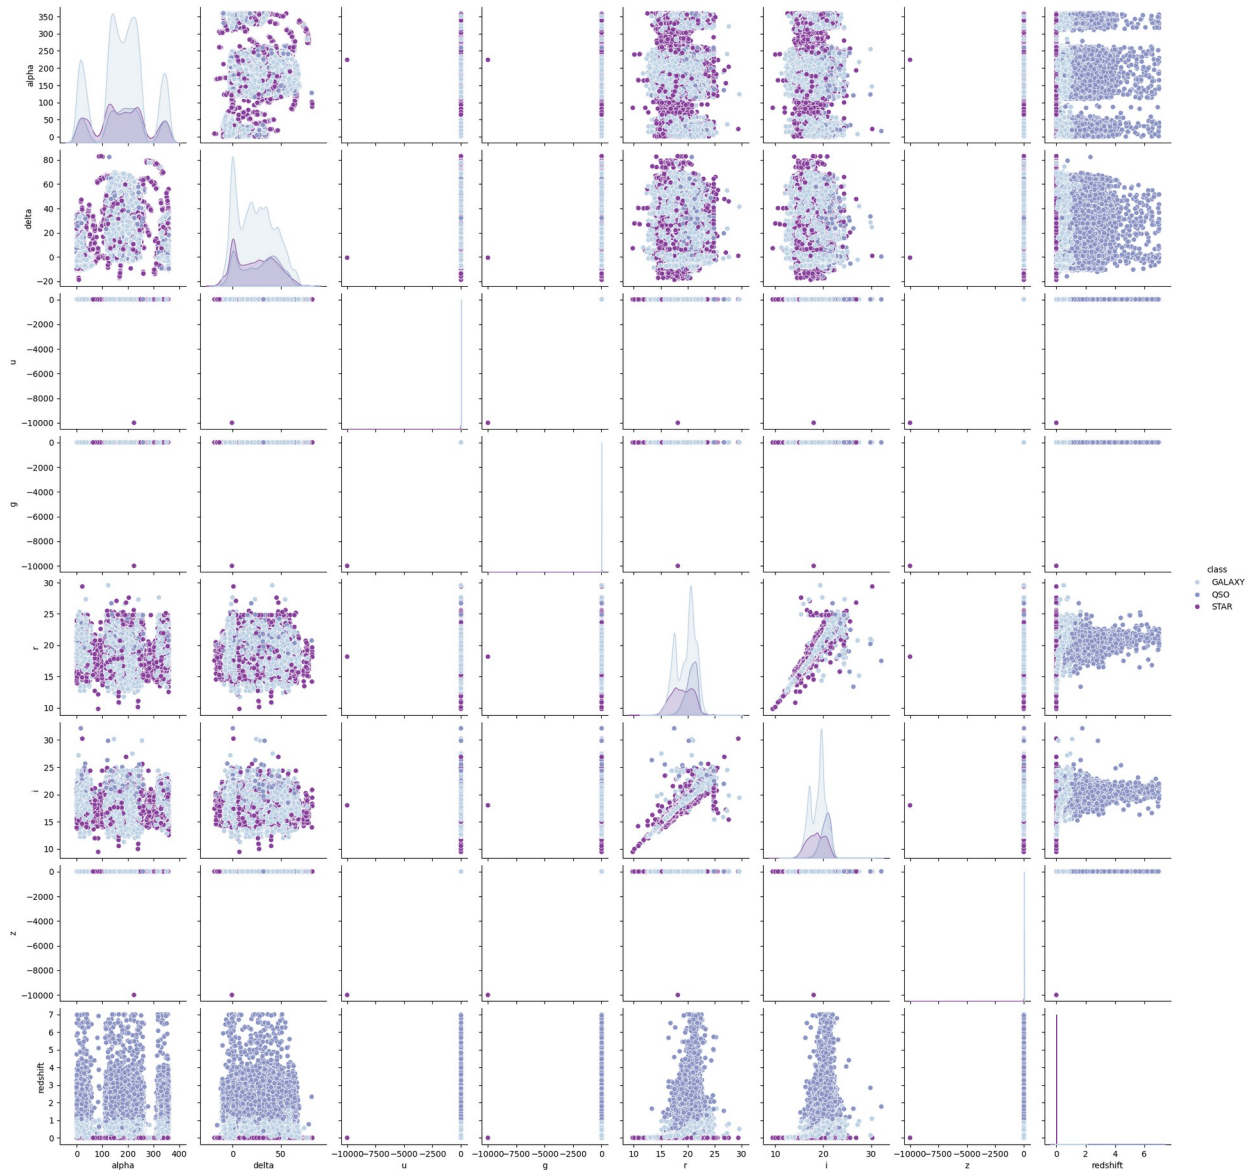
for i, col in enumerate(columns):
    sns.violinplot(x=col, data=training_data_clean, ax=axes[i],
color=colors[i])

plt.tight_layout()
plt.show()
```

Visualize the feature relationships using pairplots from seaborn

```
# make a pairplot of the training data
sns.pairplot(training_data_clean, hue = 'class', palette= 'BuPu')
plt.show()
```



We can see that the features are highly different in their distributions, and therefore should be scaled/standardized.

Data preprocessing and visualization

Split the training data into `X_train` and `y_train`

```
X_train = training_data_clean.drop(columns=['class'])
y_train = training_data_clean['class']
X_train.shape, y_train.shape

((79638, 8), (79638,))
```

Transform the categorical class variable using LabelEncoder

```
le = LabelEncoder()
# encode the target variable
y_train = le.fit_transform(y_train)
le.classes_

array(['GALAXY', 'QSO', 'STAR'], dtype=object)

le.transform(['GALAXY', 'QSO', 'STAR'])

array([0, 1, 2])
```

Look at the target variable distribution

```
(f'The distribution of the target variable is:
{np.bincount(y_train)}')
```

'The distribution of the target variable is: [47585 14802 17251]'


```
sns.countplot(x=y_train, palette='BuPu')
plt.title('Distribution of the target variable class')
plt.xlabel('Class')
```

C:\Users\anekl\AppData\Local\Temp\ipykernel_8884\2183804107.py:1:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x=y_train, palette='BuPu')
```

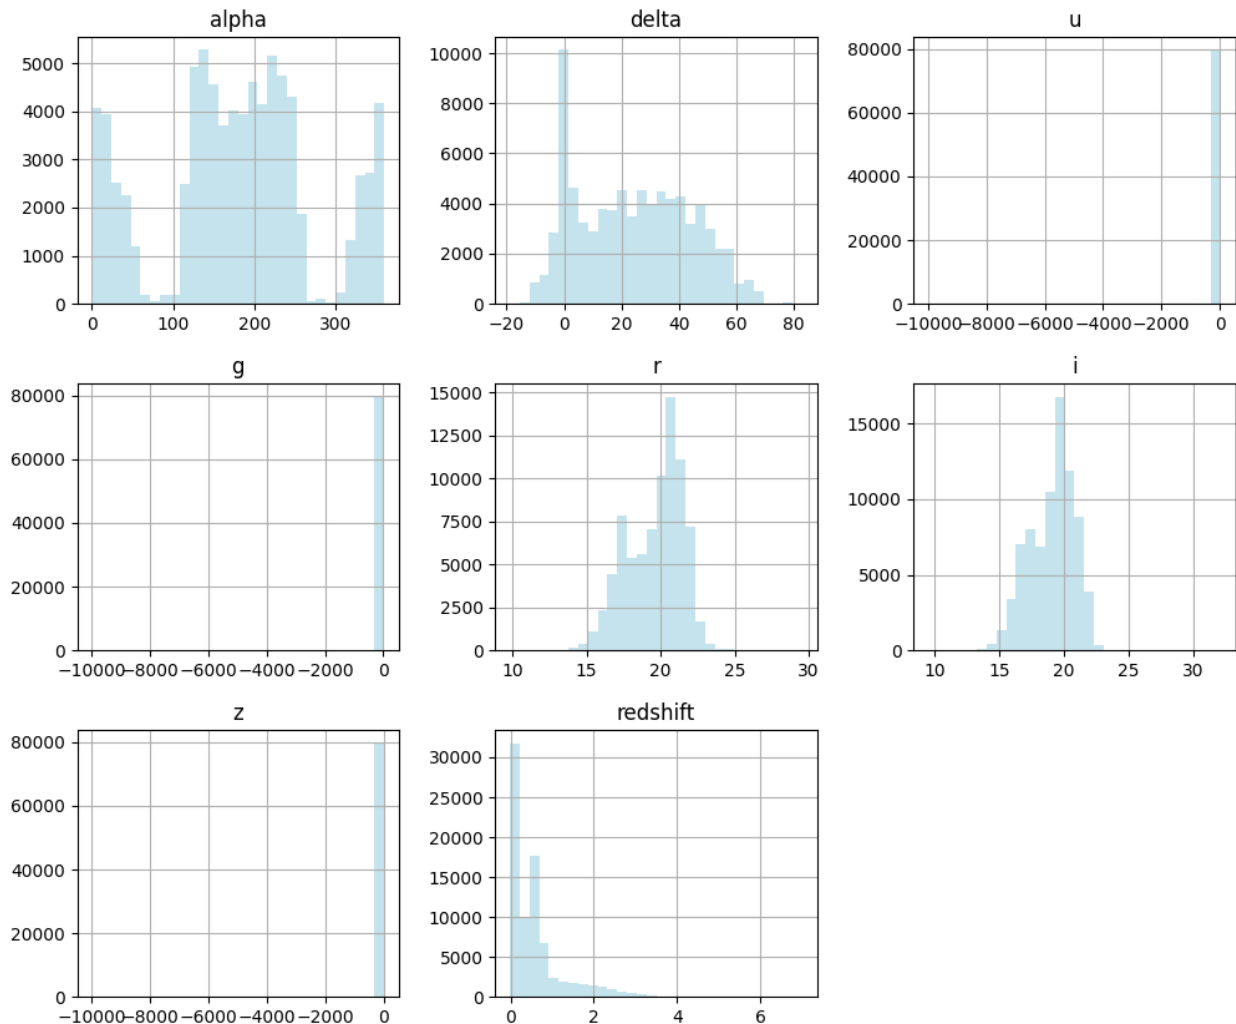
Text(0.5, 0, 'Class')



We can see from the plot above that the target variable is quite unbalanced, with more than half of the samples in the first class, galaxy.

Visualize the features before removing outliers

```
X_train.hist(figsize=(12, 10), bins=30, color='lightblue', alpha=0.7)
array([[<Axes: title={'center': 'alpha'}>,
        <Axes: title={'center': 'delta'}>, <Axes: title={'center':
'u'}>],
        [<Axes: title={'center': 'g'}>, <Axes: title={'center': 'r'}>,
        <Axes: title={'center': 'i'}>],
        [<Axes: title={'center': 'z'}>,
        <Axes: title={'center': 'redshift'}>, <Axes: >]],
dtype=object)
```



Identify and remove outliers using Z-score

```
# Detect outliers using z-scores
```

```
# Compute Z-scores
```

```
z_scores = (X_train - X_train.mean()) / X_train.std()
```

```
# Identify Outliers (absolute Z-score > 3)
```

```
outliers = (np.abs(z_scores) > 3)
```

```
# Count total outliers
```

```
num_outliers = outliers.sum().sum() # Summing over all columns
```

```
print(f"Total number of outliers: {num_outliers}")
```

```
Total number of outliers: 1831
```

```
if isinstance(y_train, np.ndarray):
```

```
    y_train = pd.Series(y_train, index=X_train.index)
```

```
# Filter rows where all features have Z-score < 3
X_train_clean = X_train[(np.abs(z_scores) < 3).all(axis=1)]

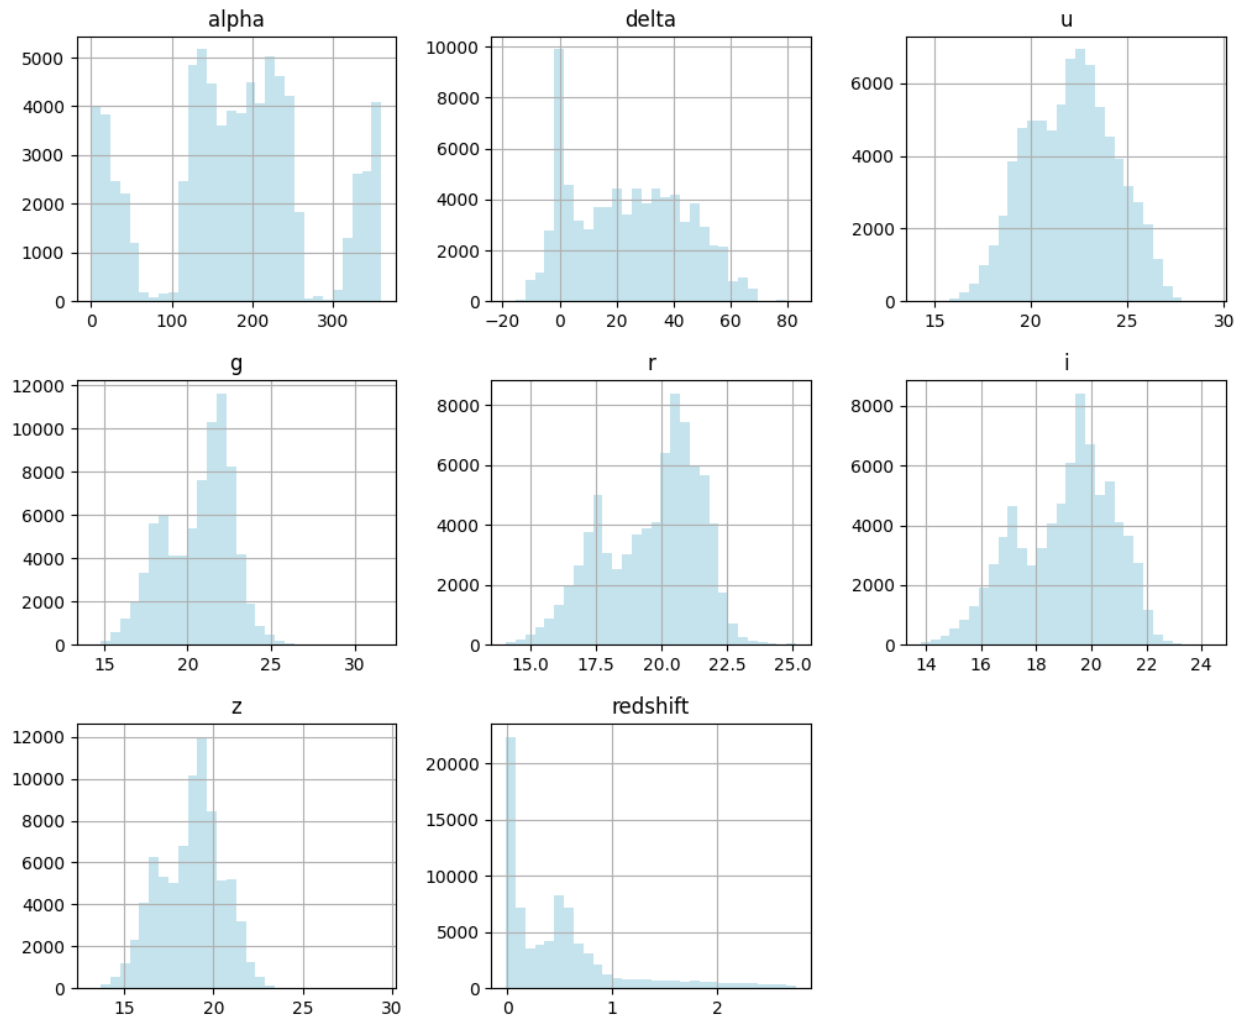
# Remove corresponding rows from y_train
y_train_clean = y_train.loc[X_train_clean.index]

# Print the shapes of the original and cleaned datasets
print("Original X_train shape:", X_train.shape)
print("Original y_train shape:", y_train.shape)
print("Cleaned X_train shape:", X_train_clean.shape)
print("Cleaned y_train shape:", y_train_clean.shape)

Original X_train shape: (79638, 8)
Original y_train shape: (79638,)
Cleaned X_train shape: (77943, 8)
Cleaned y_train shape: (77943,)
```

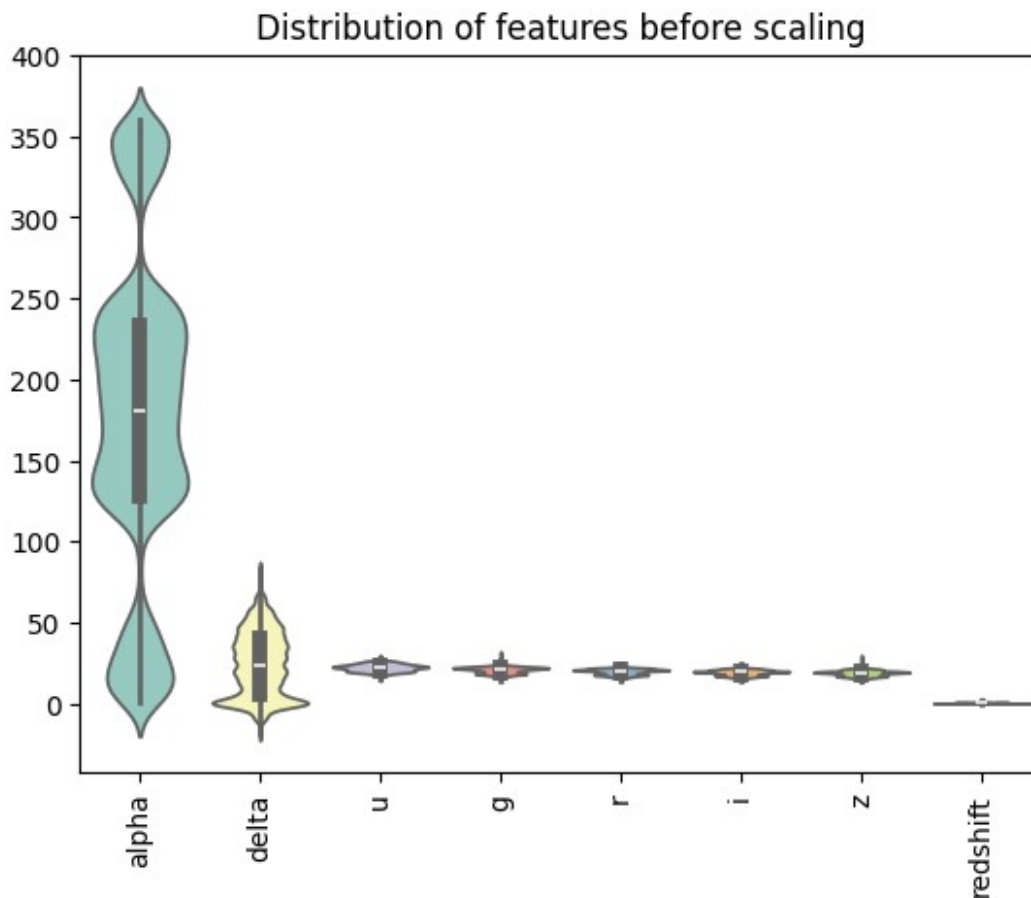
Visualize the features after removing outliers

```
histogram = X_train_clean.hist(figsize=(12, 10), bins=30,
color='lightblue', alpha=0.7)
```



Visualize the data before scaling

```
sns.violinplot(X_train_clean, palette = sns.color_palette('Set3',
n_colors= X_train_clean.shape[1]))
plt.xticks(rotation = 90)
plt.title('Distribution of features before scaling')
plt.show()
```



Modelling

Define pipelines

SVM

```
# Define the SVC pipeline:
pipe_svc = make_pipeline(StandardScaler(), SVC(random_state=1))
pipe_svc.get_params()

{'memory': None,
 'steps': [('standardscaler', StandardScaler()), ('svc',
SVC(random_state=1))],
 'verbose': False,
 'standardscaler': StandardScaler(),
 'svc': SVC(random_state=1),
 'standardscaler__copy': True,
 'standardscaler__with_mean': True,
 'standardscaler__with_std': True,
 'svc__C': 1.0,
 'svc__break_ties': False,
 'svc__cache_size': 200,
 'svc__class_weight': None,
```



```

'svc__coef0': 0.0,
'svc__decision_function_shape': 'ovr',
'svc__degree': 3,
'svc__gamma': 'scale',
'svc__kernel': 'rbf',
'svc__max_iter': -1,
'svc__probability': False,
'svc__random_state': 1,
'svc__shrinking': True,
'svc__tol': 0.001,
'svc__verbose': False}

```

Logistic regression

```

# Define the logistic regression pipeline:
pipe_logistic = make_pipeline(StandardScaler(), PCA(n_components=7),
LogisticRegression(random_state=1))
pipe_logistic.get_params()

{'memory': None,
 'steps': [('standardscaler', StandardScaler()),
 ('pca', PCA(n_components=7)),
 ('logisticregression', LogisticRegression(random_state=1))],
 'verbose': False,
 'standardscaler': StandardScaler(),
 'pca': PCA(n_components=7),
 'logisticregression': LogisticRegression(random_state=1),
 'standardscaler__copy': True,
 'standardscaler__with_mean': True,
 'standardscaler__with_std': True,
 'pca__copy': True,
 'pca__iterated_power': 'auto',
 'pca__n_components': 7,
 'pca__n_oversamples': 10,
 'pca__power_iteration_normalizer': 'auto',
 'pca__random_state': None,
 'pca__svd_solver': 'auto',
 'pca__tol': 0.0,
 'pca__whiten': False,
 'logisticregression__C': 1.0,
 'logisticregression__class_weight': None,
 'logisticregression__dual': False,
 'logisticregression__fit_intercept': True,
 'logisticregression__intercept_scaling': 1,
 'logisticregression__l1_ratio': None,
 'logisticregression__max_iter': 100,
 'logisticregression__multi_class': 'auto',
 'logisticregression__n_jobs': None,
 'logisticregression__penalty': 'l2',
 'logisticregression__random_state': 1,

```

```
'logisticregression__solver': 'lbfgs',
'logisticregression__tol': 0.0001,
'logisticregression__verbose': 0,
'logisticregression__warm_start': False}
```

Random Forest

```
# Define the random forest pipeline:
pipe_rf = make_pipeline(RandomForestClassifier(random_state=1))
pipe_rf.get_params()

{'memory': None,
 'steps': [('randomforestclassifier',
RandomForestClassifier(random_state=1))],
 'verbose': False,
 'randomforestclassifier': RandomForestClassifier(random_state=1),
 'randomforestclassifier__bootstrap': True,
 'randomforestclassifier__ccp_alpha': 0.0,
 'randomforestclassifier__class_weight': None,
 'randomforestclassifier__criterion': 'gini',
 'randomforestclassifier__max_depth': None,
 'randomforestclassifier__max_features': 'sqrt',
 'randomforestclassifier__max_leaf_nodes': None,
 'randomforestclassifier__max_samples': None,
 'randomforestclassifier__min_impurity_decrease': 0.0,
 'randomforestclassifier__min_samples_leaf': 1,
 'randomforestclassifier__min_samples_split': 2,
 'randomforestclassifier__min_weight_fraction_leaf': 0.0,
 'randomforestclassifier__monotonic_cst': None,
 'randomforestclassifier__n_estimators': 100,
 'randomforestclassifier__n_jobs': None,
 'randomforestclassifier__oob_score': False,
 'randomforestclassifier__random_state': 1,
 'randomforestclassifier__verbose': 0,
 'randomforestclassifier__warm_start': False}
```

K-nearest-neighbour (KNN)

```
# Define the KNN pipeline:
pipe_knn = make_pipeline(StandardScaler(), KNeighborsClassifier())
pipe_knn.get_params()

{'memory': None,
 'steps': [('standardscaler', StandardScaler()),
 ('kneighborsclassifier', KNeighborsClassifier())],
 'verbose': False,
 'standardscaler': StandardScaler(),
 'kneighborsclassifier': KNeighborsClassifier(),
 'standardscaler__copy': True,
 'standardscaler__with_mean': True,
 'standardscaler__with_std': True,
```

```
'kneighborsclassifier__algorithm': 'auto',
'kneighborsclassifier__leaf_size': 30,
'kneighborsclassifier__metric': 'minkowski',
'kneighborsclassifier__metric_params': None,
'kneighborsclassifier__n_jobs': None,
'kneighborsclassifier__n_neighbors': 5,
'kneighborsclassifier__p': 2,
'kneighborsclassifier__weights': 'uniform'}
```

Evaluate different models and hyperparameters using GridSearchCV with cross-validation

SVM

```
cv_strategy = StratifiedKFold(n_splits=5, shuffle=True,
random_state=42)

gs_svm = GridSearchCV(estimator=pipe_svc,
                      param_grid={'svc__C': [0.1, 1, 10, 100],
'svc__kernel': ['rbf']},
                      scoring='f1_macro',
                      cv=cv_strategy,
                      n_jobs=-1)

gs_svm = gs_svm.fit(X_train_clean, y_train_clean)
print(gs_svm.best_score_)
print(gs_svm.best_params_)

best_model_svm = gs_svm.best_estimator_

0.9666694954815223
{'svc__C': 100, 'svc__kernel': 'rbf'}
```

Logistic regression

```
cv_strategy = StratifiedKFold(n_splits=5, shuffle=True,
random_state=42)

gs_logistic = GridSearchCV(estimator=pipe_logistic,
                           param_grid={'logisticregression__penalty': ['l2'],
'logisticregression__C': [1,10,100]},
                           scoring='f1_macro',
                           cv=cv_strategy,
                           n_jobs=-1)

gs_logistic = gs_logistic.fit(X_train_clean, y_train_clean)
print(gs_logistic.best_score_)
print(gs_logistic.best_params_)

best_model_logistic = gs_logistic.best_estimator_
```

```
0.9529143717852919
{'logisticregression__C': 100, 'logisticregression__penalty': 'l2'}
```

Random forest

```
gs_rf = GridSearchCV(estimator=pipe_rf,
                     param_grid={'randomforestclassifier__n_estimators':
[80, 100, 120],
                                'randomforestclassifier__max_features':
['sqrt', 'log2'],
                                'randomforestclassifier__max_depth':
[5, 10, 15],
                                'randomforestclassifier__criterion':
['entropy']},
                     scoring='f1_macro',
                     cv=5,
                     n_jobs=-1)

gs_rf = gs_rf.fit(X_train_clean, y_train_clean)
print(gs_rf.best_score_)
print(gs_rf.best_params_)

best_model_rf = gs_rf.best_estimator_

0.9730246117133078
{'randomforestclassifier__criterion': 'entropy',
 'randomforestclassifier__max_depth': 15,
 'randomforestclassifier__max_features': 'log2',
 'randomforestclassifier__n_estimators': 80}
```

KNN

```
cv_strategy = StratifiedKFold(n_splits=5, shuffle=True,
                               random_state=42)

gs_knn = GridSearchCV(estimator=pipe_knn,
                     param_grid={'kneighborsclassifier__n_neighbors': [3,
5, 7],
                                'kneighborsclassifier__weights':
['distance'],
                                'kneighborsclassifier__metric':
['euclidean', 'manhattan']},
                     scoring='f1_macro',
                     cv=cv_strategy,
                     n_jobs=-1)

gs_knn = gs_knn.fit(X_train_clean, y_train_clean)
print(gs_knn.best_score_)
print(gs_knn.best_params_)

best_model_knn = gs_knn.best_estimator_
```

```
0.9404996981321396
{'kneighborsclassifier__metric': 'manhattan',
 'kneighborsclassifier__n_neighbors': 5,
 'kneighborsclassifier__weights': 'distance'}
```

Build classifier based on all training samples using the "optimal parameters"

The code below is based on the best model of the Random Forest Classifier, since this had the highest score during GridSearch.

All fitting, predicting and evaluation is done using this classifier.

```
# fit the best model to the training data (random forest)
best_model_rf.fit(X_train_clean, y_train_clean)

Pipeline(steps=[('randomforestclassifier',
                  RandomForestClassifier(criterion='entropy',
max_depth=15,
n_estimators=80,
max_features='log2',
random_state=1))])
```

Evaluate model performance

Confusion matrix

```
# confusion matrix for the random forest model using train/test split
on training data
X_train_data, X_test_data, y_train_labels, y_testlabels =
train_test_split(X_train_clean, y_train_clean, test_size=0.4,
random_state=42)
y_pred = best_model_rf.predict(X_test_data)

confusion = confusion_matrix(y_testlabels, y_pred)
confusion

array([[18771,    38,     1],
       [  160,  5295,     0],
       [     0,     0,  6913]])
```


Plot the confusion matrix using matplotlib

```
# Code borrowed from lecture 'Chapter_6_part_2b'

fig, ax = plt.subplots(figsize=(3, 3))
ax.matshow(confusion, cmap=plt.cm.Blues, alpha=0.3)
for i in range(confusion.shape[0]):
    for j in range(confusion.shape[1]):
        ax.text(x=j, y=i, s=confusion[i, j], va='center', ha='center')

plt.xlabel('Predicted label')
```

```
plt.ylabel('True label')
plt.tight_layout()
plt.show()
```



Model performance using classification report in Scikit learn

```
classification_rep = pd.DataFrame(classification_report(y_testlabels,
y_pred, output_dict=True)).T
classification_rep
```

	precision	recall	f1-score	support
0	0.991548	0.997927	0.994727	18810.000000
1	0.992875	0.970669	0.981646	5455.000000
2	0.999855	1.000000	0.999928	6913.000000
accuracy	0.993617	0.993617	0.993617	0.993617
macro avg	0.994759	0.989532	0.992100	31178.000000
weighted avg	0.993622	0.993617	0.993592	31178.000000

We can see that the output from the classification report gave the same F1 scores as we manually calculated above.

Kaggle submission

```
y_test = best_model_rf.predict(test_data_clean)
y_test = pd.DataFrame(y_test, columns=["class"])
y_test.index.name = "ID"
y_test[['class']].to_csv("data/sample_submission.csv")
```